Scare vs non-scare quotes: Annotation guidelines

Aina Garí Soler
September 4th 2025

Table of contents

Table of contents 1
1. Introduction 2
2. Annotation labels 2
2.1 High-level annotation 3
Non-scare quotes (NSQ) 3
Scare quotes (SQ) 3
Ambiguous 4
Unsure 4
2.2 Low-level annotation 4
2.2.1 Types of non-scare quotes 5
● Meta-linguistic comment 5
● Direct quoting 5
● Mixed quoting without attitude 7
● Emphasis 7
● Designation 7
● Ambiguous 8
2.2.2 Types of scare quotes 8
● SQ Usage 9
● SQ Word 10
● Mixed quotation with attitude 10
● FLAG: Reusing a word previously used by another speaker in the conversation 11
● Ambiguous 13
2.3 Common types of double annotations and ambiguities 13
● Meta-linguistic comment + SQ usage/word 13
● Ambiguity: Mixed quotation with/without attitude 14
● SQ usage + Designation 14
● Meta-linguistic comment + Designation 14
● Meta-linguistic comment + Direct quoting 14
3. Annotation Platform 15
3.1 Label Studio basic usage 15

3.2 Reddit and format specificities 17
Quotations 17
Edited messages 17
Reddit / winning-args corpus terminology 18
4. Simplified action plan 18

1. Introduction

This document is intended to serve as training and as a reference for annotators to maximize consistency in the annotation of usages of quotation marks, as well as for users of the corpus who want to gain a deeper understanding of the meaning of each label. This is the version of the guidelines that was provided to the second annotator when annotating the corpus. For a version of the guidelines containing more examples, visit the project's repository.

Quotation marks (QM, “, ‘, «) are a feature of written discourse and can serve multiple purposes. In this project, we are annotating QM usages (QMUs) of up to three words found in social media forums. We are particularly interested in identifying cases of “**scare quoting,”** which express some sort of attitude or distancing toward the quoted word (more information in Section Scare quotes).

In section 2, we present each annotation label in detail. Section 3 contains information on the annotation platform and how it works, as well as formatting specificities of the corpus that has been selected for annotation. Section 4 contains a brief action plan summarizing the decisions to make when annotating.

2. Annotation labels

Annotation takes place at two hierarchical levels, as illustrated in the tree below:

- **High-level annotation**: first, a decision needs to be made on whether a specific quoted usage is an instance of scare quoting (SQ) or not (NSQ).
- **Fine-grained annotation**: then, the specific function of the quotation marks needs to be selected from a more fine-grained set of labels.

The annotator **must provide at least one label at each annotation level**. The high-level and fine-grained labels must be **compatible**: for example, following the tree below, annotating an instance with only “SQ” and “Emphasis” (which is under NSQ) will be considered invalid.

It is perfectly valid and even sometimes necessary to annotate multiple labels per level for a single QMU. Rather than a classification task, this annotation can be seen as a *characterization* task, and labels can in fact be interpreted as *traits* of a QMU.

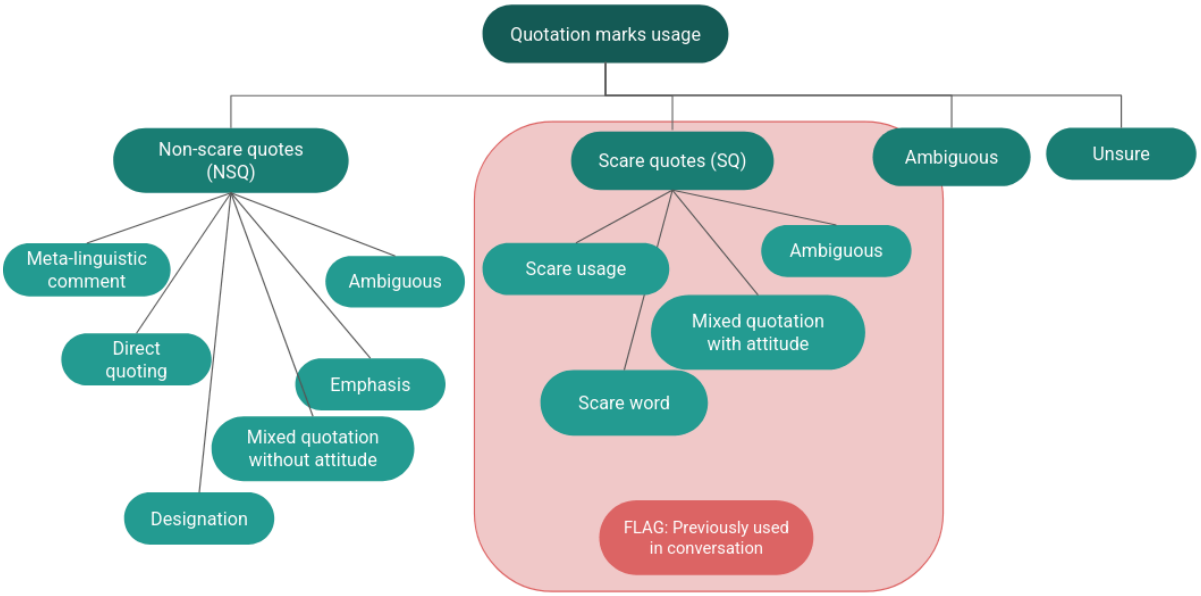

2.1 High-level annotation

Here we provide a general description of each high-level class. Section 2.2 elaborates on the subtypes of each type of quoting, providing real examples which can further clarify this distinction and facilitate the decision of annotation at this level.

There are 4 labels in this category: Non-scare quotes (NSQ), Scare quotes (SQ), Ambiguous and Unsure. **Note that you can select more than one label for a single instance.**

Non-scare quotes (NSQ)

The main characteristic of non-scare quotes is that the quoted passages are to be taken and interpreted at face value. There is no added attitude, no intention by the writer to warn the reader about a non-standard word or word usage or to distance themselves from what is said. The two most well-known functions of non-scare quotes are **meta-linguistic mentioning** (*"Cat" has three letters*) and **citing** someone else's words (*He said "I'm hungry"*). They can also be used to mark specialized **terminology or titles** (*Have you watched "Gone with the wind?"*). More details can be found in Section Types of non-scare quotes.

Scare quotes (SQ)

We consider a quoted word or phrase to be an instance of scare quotes if one of the following is true:

- The word is not to be taken at face value. Perhaps because it is a suboptimal word choice (e.g., because it doesn't fully apply to the situation being described, or because the writer is unsure that this is the correct word to use), because it is used in a figurative sense, or with irony (e.g., *the "generous" lady donated one cent to charity*)
- The writer seems to want to distance themselves from the usage. Perhaps because the word is connotatively loaded or controversial, or because it comes from a different register (e.g., *This lady works as an "influencer" on Instagram)*.
- The word itself is very unusual, incorrect, or made-up.
- The writer is quoting someone else's words, but expressing, implicitly or explicitly, some attitude towards what is said (for example, that they would have said it in a different way): *he said he was "the humblest person".*

Ambiguous

Double annotation vs ambiguity
Sometimes, a QMU has elements of both SQ and NSQ simultaneously, which are not in contradiction with each other. In this case, the annotator should simply select both categories in the annotation interface. We will refer to this as a **double annotation**.
However, it is possible that you see elements of both SQ and NSQ but a reading of type SQ would not be compatible with a simultaneous NSQ reading. These can be cases where we lack context to determine with certainty the exact category of the QMU. In this case, an instance should be annotated as **Ambiguous**.

Unsure

If the annotator is unsure of their choice, the item can additionally be flagged with the label "Unsure". This is useful to mark difficult cases, either to a) revisit them later for further consideration; or b) keep track of hard cases to enrich a later analysis.

The annotator is free to add this label to any difficult cases and to keep it there even after they have made a final decision for an instance, to leave a trace of their doubts. However, **at the end of the annotation process, no instance annotated with only "Unsure" should be left**: a decision needs to be made for every instance.

2.2 Low-level annotation

In what follows, we describe in detail every subtype of non-scare and scare quotes with multiple examples.

Note that it is possible, for a given item, to have a single higher-level annotation (regarding whether it is SQ or NSQ) but to have a double annotation at this lower level. We also discuss typical cases of double annotation in Section 2.3.

2.2.1 Types of non-scare quotes

- Meta-linguistic comment

This label is used for meta-linguistic usages of a word or phrase, also known as *mentions* (rather than *uses*) of a word. These are cases where the word is talked about “as a word” and is not used with its existing meanings. The meta-linguistic comment may make reference to the form or sound of the word, but also to its meaning. Meta-linguistic comments can define, discuss or even challenge the meaning of a word.

One way to identify many, but not necessarily all of such cases is to imagine whether the QMU could naturally be preceded by “the word”, “the term” or “the phrase”, which would be referring to what is in quotes. However, there can also be meta-linguistic comments of (short) sentences or statements.

fake450_##_t1_cue9cr1_##_rt-t3_3i8nmh: This is my first time encountering this quote, but tis is my interpretation: This quote could be held as true only under the premise that "**religion**" in this case does not refer exclusively to theistic religions. Any system of rituals or beliefs that people follow, including the patriotic following of orders you referred to, is a religion in this context. As for your other examples, they might be countered in a small variety of ways, depending on the interpreter's beliefs, including something like the following: -Actions taken by a person in poverty in order to survive are not evil because the person has no alternative -People who are drugged are not themselves, in that they are controlled by the effects of the substance or addiction thereto.

ID 7194

Finally, when a short definition is given for a word and is put between quotes, we consider this as a meta-linguistic comment as well:

"Terrible" used to mean "inspiring terror," which was a good thing. "Awful" used to mean "**inspiring awe,**" which was a good thing. Same with "Fearful." Which is why the King of

ID 5725

- Direct quoting

This label is used for instances where the content inside quotes is to be understood as something that has been uttered (said, written, typed, thought, googled) by someone else. Contrary to the case of mixed quotation introduced below, here the utterance is reproduced in full and in direct speech (*he said "I'm hungry"*), and is not embedded in the writer's own words (*he said he was "hungry"*).

bazmeg_##_t1_cm5xgmh_##_rt-t3_2mikc0: Serious, clinical depression gets noticed. Lighter depression is a philosophical issue. Are you supposed to be happy and joyful if your life is objectively not satisfying enough? Should everybody regardless of their conditions take a happy pill and run around smiling at people and answer "**great!**" if they ask how are you?

ID 597

Note that the content between QM need not have been uttered in reality by someone. It can be an **imaginary** utterance, or a hypothetical. And given what has been observed in the data, we adopt a flexible definition of direct quoting which does not require the presence of a reporting verb such as "say" or "ask," as long as it is clear from the context that the words in quotes are cited from someone else. These are cases where we can imagine a phrase like "X said" potentially being appended to it:

space, time, and matter." my question is, how exactly is it that you are proposing the right to firearms derives from nature? "**It just does**" is not a sufficient answer. https://en.m.wikipedia.org/wiki/Natural_and_legal_rights Explain why we

ID 6953

caw81_##_t1_ceee0tu_##_rt-t1_ceebjb6: http://en.wikipedia.org/wiki/Poverty_in_the_United_States > In November 2012 the U.S. Census Bureau said more than 16% of the population lived in poverty in the United States,

Valridagan_##_t1_ceee8bh_##_rt-t1_ceee0tu: Alright, thank you. I apologize again. ...Wait. "**More than**"? How much more?

ID 3962

However, **we do not count as instances of direct quoting** meta-linguistic comments about how people use a word in general or often (e.g., *when people talk about X…, they mean…*; *people would call this X, but…*). These are annotated as **meta-linguistic comments** instead:

discussion is happening to begin with. [END-CITE]That's not what people mean when they are talking about a "**semantic argument.**" When people use that term, they mean that rather than debating the merits of the point that the person was making, they are instead fixating on how it

ID 3506

- Mixed quoting without attitude

Mixed quoting is similar to direct quoting, but the content inside the quotes is embedded in the writer's sentence and is often not a full utterance. This sort of selective quoting can be done as a way to express some attitude towards what the other person said, but in this case this would be an instance of SQ. Mixed quotation **without** attitude is reserved for cases where there is no attitude or judgment.

of punching your friend remains non-zero [END-CITE]I didn't say anything about the possibility. I said I have the "**ability**" to do it. I used that word because that's the word you

ID 6778

- Emphasis

Although it is sometimes considered to be an incorrect use of QM, they are sometimes used simply as a means to emphasize a word in a similar way as it could be done with boldface, italics or uppercase.
The QM draw attention to a word or phrase (in the example below, to emphasize the contrast with "under") without questioning its appropriateness.

swearrengen_##_t1_cp3l5xc_##_rt-t1_cp3jczv: But Hydropower can change how much goes "**over**" the waterfall versus "under" (via tunnels), so stratys3 may be right on the point he was trying to make?

ID 6042

- Designation

This is a broad umbrella category for the following types of phenomena:

a) **terminology**: using specialized terminology or concepts that may not be known to everyone;

BaneFlare_##_t1_ct6qul7_##_rt-t1_ct6kvem: CS:GO seems to suffer from the "**garage tournament**" effect even more often that Dota 2 in my experience.

ID 8210

b) **titles & artifacts**: delimiting proper names of works (e.g. book titles), objects, or more abstract ideas:

Ferrousity_##_t1_cumlhj5_##_rt-t3_3j348p: But these games usually come with a way to circumvent it. FF has items with the "**No Encounter**" effect, and pokemon has Repels (or just run if your speed/AGL stat is high enough). I get that it can be problematic, but I can't imagine how else you can really do non-random encounters. Later installments of FF (XII, XII+) have an ATB (Active Time Battle)

ID 6831

c) **Objectification/Conceptualization**: using words as conceptual objects or, mostly in the case of phrases, as lexicalized units, potentially acknowledging that the sentence would be ungrammatical if the phrase is not understood as a unit. This also includes using a word in an unusual part of speech or in a somewhat ungrammatical way.

IWriteWithATalon_##_t1_cu0tiej_##_rt-t1_cu0t7ty: I kind of addressed this in my post; you're not offering an alternative form of work. You're just saying "You can't do this job". Not "Okay, do this one instead". There's no choice for a minimum wage job; it's just a choice of "**Prostitute**" or "Starve". Yeah, it sucks, but if you take away Option A, all that's left is a slow death by Option B.

ID 7496

paneubert_##_t1_cmj9bg4_##_rt-t1_cmj8xrp: Agree about the time line, but there could be many other substances that would work in much the same way that were present at that time. Not that I am subscribing to the "**he drugged her**" theory. Just playing devils advocate.

ID 881

those who believe in a downward causation are concerned with the question of 'why'. The tools of science aren't equipped to answer the '**why**' just as the tools of philosophy are inadequate to answer 'how'. They seem dogmatic when a) someone thinks they are exclusive ways to view

ID 9167

This last example could be interpreted as SQ in that a word may be used in an incorrect way, but we consider it a case of NSQ because there does not seem to be anything unusual going on semantically and there is nothing unusual in the use of the word.

- Ambiguous

To be used on top of the selected labels if there is an ambiguity between two or more NSQ labels. It indicates that the selected labels correspond to different, incompatible readings.

2.2.2 Types of scare quotes

The fine-grained categories we have for SQ are more coarse-grained than those for NSQ, due to the difficulty in determining the specific reason why someone decided to use them. In summary, we have three categories:

- SQ usage (an umbrella category for most SQ)
- SQ word (for cases when the word in quotes is made-up or incorrect)
- Mixed quotation with attitude

Additionally, we have an optional flag to mark cases where the passage in quotes has previously been used in the conversation. This is a useful way to mark, in dialog, QMU where the writer reuses (and perhaps contests) someone else's word choice.

- SQ Usage

As explained above, this label is a general umbrella category for any SQ which does not fall into the other two categories. It marks any non-standard, suboptimal, figurative or dubious word usage, except for those that stem from the word form itself (which would fall under "SQ word") or which are mixed quotations with attitude.
It is often hard to determine the exact reason or type of scare usage, but this label should be annotated if one or more of the following are perceived in the usage of a word or phrase:

- ironic uses of a word;
- or figurative uses of a word;
- slang or words that do not match the register of the conversation;
- controversial words or concepts;
- expressing that a word's meaning is vague, unclear or ambiguous, particularly if in a negative light;
- words that do not apply to a specific situation;
- word usages that are inaccurate or inexact in a given context;
- the writer wishes to distance themselves from the word for some reason.

Example of marking a controversial phrase:

Daksund_##_t1_cj02zax_##_rt-t1_cj02sy5: Again, eugenics is less about destroying '**bad genes**' and more about endorsing 'good genes'. We can see clearly that some people are more

ID 2726

A figurative usage:

mastermind42_##_t1_ccm0kcp_##_rt-t3_1ntobs: I am going to just address the suicide argument since everything else has been already critiqued. ONE: Suicide is a waste of human potential. It is important is to get more out of the investment of time and resources that go into raising a person then it cost to raise them. In order to raise someone up to the age of (for example) 18, that is 18 years in which this person consumed resources, time, and effort. If it lived in a first world country, this included teachers, doctors, massive amount of food, nonrenewable resources, etc. It consumed ALL of this on a implied social debt to society. It "**borrowed**" from the bank of humanity with the promise of at least trying to repay it. Whatever

ID 4319

Sometimes, a SQ usage comes accompanied by a negation or some other negative expression which itself expresses that the word does not apply to a specific situation or that the writer does not approve of the word. While one might consider that the scare function is already done outside of the QMU, we annotate these as SQ usages. In the example below, this negative appreciation is done through “I just disagree”:

> advanced nation. I don't disagree with that, I just disagree that Germany was a particularly "**civilized**" nation. What kind of civilized nation would do a genocide? I guess you're getting close to changing my

ID 6629

Similarly, in the case below, the person expresses that the word “married” does not really apply to this situation, as in their opinion, such a marriage would be “void”:

> forms as a married person. So, in that case, I'm not really mattered in any sense. If the only way that a person can get "**married**" is to do so in a way that gives none of the rights that other groups of people who are married get is exactly equal. Their marriages are void. They

ID 596

- SQ Word

This label serves to annotate SQ usages whose non-standardness stems from the word form itself: the word may be made-up, incorrect, or so unusual that its correctness is doubted upon.
Contrary to SQ Usage, this type of SQ does not have to do with the meaning of the word itself or the sense in which it is used in the given context.
Words from other languages (code-switching) also belong to this category. This category mostly applies to words rather than phrases.

> **SidewaysGate_##_t1_ce448ho_##_rt-t1_ce442e3**: I'm confused. Do you want it to be more about the physical computers that we're using? My program focused a fair amount of that. For example there was a pretty significant section early on explaining the effect of locality on program efficiency and an explanation of how the hardware actually matters while programming. We also talked about some older techniques and did some assembly programming. We took measurements in profiling our software and did write-ups on it. This sounds pretty '**science-y**' to me.

ID 4095

- Mixed quotation with attitude

Mixed quotation is used to cite <u>someone else</u>’s words, but the content inside quotes is embedded in the writer’s sentence and is often not a full utterance.

This class is similar to the NSQ class of “Mixed quotation without attitude,” but in this case there is a clear attitude or judgment towards what the other person said.
Note that, as with the mixed and direct quoting classes in NSQ, we consider utterances that may not have been said in reality, but are imaginary or hypothetical.
In Reddit data, it is often the case that people **reuse other users**’ words in QM to indicate their disapproval at the word choice. Such cases do **not necessarily** constitute examples of mixed quotation with attitude, and have their own flag, which is the next label described in this section. **The label “Mixed quotation with attitude” should only be used if it is clearly and explicitly expressed, in the writer’s text, that the expression between quotes is (or may be) taken from something someone has said/written/uttered.**

LiveBeef_##_t3_306jm9: It is not uncommon to hear in sports press conferences of a player whose team is Super Bowl-bound, playoff-bound, or otherwise enjoying great success that they are "**humbled**" for the opportunity to play at that level. Or a winning candidate from an election to be "humbled" by the turnout in their favor. The problem is, when they say that, what they are really feeling is great pride and sense of accomplishment, which is decisively not humility. "Humbled" often goes hand-in-hand with "Humiliated", which is the opposite of what most people who describe themselves as "humbled" have gone through. If someone wants to say "Although I

ID 5705

blocking software becomes, the less value online advertising will become. You'll see websites charging for service or offering lower quality as a result of reduced income from advertisers. Also you are saying that they "**should**" alloy you to not see ads. Why should they do that? Obviously that would benefit you, but what do the websites have to gain from that?

ID 1805

Note that for things that have been said within the same Reddit conversation, but which are not expressed in the form of a mixed quotation (embedded in the writer’s sentence), the label presented next (FLAG) may be more appropriate.

Finally, note that in all kinds of citation (direct and mixed quotation, with and without attitude), it is possible to cite oneself. In such cases, mixed quotation with attitude may look like someone who changed their mind or who cites a past self they do not identify with anymore.

- FLAG: Reusing a word previously used by another speaker in the conversation

This flag is used to mark SQ cases where the writer is reusing a word that **someone else** in the same thread has **already** used. The writer puts the word or phrase in quotes to express their disapproval at that specific word choice or to distance themselves from it.

This is not a subclass of SQ per se, but it is a flag that can be added in addition to any other SQ label. The purposes of this flag are 1) to mark words where there is a potential disagreement between conversation participants, 2) to better characterize the reason for

using QM, and ultimately 3) to better understand the dynamic and interpersonal nature of QMUs in dialog.

This flag is different from Mixed quotation with attitude in that:

1) The person who originally wrote the quoted word or passage has participated in the thread,
2) it is not necessarily said explicitly that this word is being quoted from someone else,
3) The quoted material needs to come from a **different user**.

To add this flag, it is necessary to look at the preceding posts and see whether that word (or a very similar one) has been used by another user.

To assist in this process and facilitate the annotation of this flag, in the annotation platform, the column "**used_previously_in_subthread**" indicates whether the word has already been used in the preceding messages of the same subthread, which can facilitate annotation for this flag. However, note that the value of this column does not fully determine whether an instance deserves the flag. For example:

- Even if the column displays FALSE, a speaker may have used a slightly different word form (often sharing morphemes), or might have written with a typo (e.g., *happy*/*unhappy*, *marry/merry*).
- Even if the column displays TRUE, maybe the word has been used in an entirely different context or with a different sense, or by the QM writer themselves only. The annotator should always check the conversational context to decide.

The person who introduced the word in the conversation for the first time may be the same as the writer of the QM that is being annotated, even if other people subsequently used the same term. In this case, the QMU at hand is less likely to merit this flag. The flag should only be used if the annotator suspects that the QM are somehow criticizing another user's word usage.

MooseAtWork_##_t1_ces5dj5_##_rt-t1_ces50fs: [STA-CITE]>Photographers were reviled by painters, digital photographers reviled by film photographers, photoshoppers reviled by "pure digital photographers," and so on and so forth. [END-CITE]I'm not "**reviling**" digital cinematography at all. However, I am recognizing the two are very different beasts in the skill it takes to execute either and, in recognition of that, they need to be judged separately. You're

ID 3796

Or in this interaction:

TwizzlesMcNasty_##_t1_cc772pa_##_rt-t1_cc76puq: Thank you for responding. I think that a complaint in that situation is unwarranted because they have been done a favor by being excluded. I was thinking about this a remembered something that happened to me about 10 years back. I lived with a couple of other guys and we all hung out. One of the others got married and asked the other to be a groomsman. I was not invited. I did not want to go but I relished the chance to proclaim that I was never invited. Why would I do this. It did not really bother me and if invited I would have tried to come up with an excuse. I think I had to be enjoying playing the victim which is not a healthy attitude to have. I understand what you are saying but I still do not think it is a reason to complain.

mikalaranda_##_t1_cc77fzc_##_rt-t1_cc772pa: I think the intentions behind your complaints can make all the difference. Complaining simply because you feel like it, i.e. venting, is most likely unproductive and could be considered an "**unhealthy**" practice, as you state, for the individual engaging in the act. Complaining because you want to express to someone that they made a mistake, breached common etiquette, etc., and that they should try to be more considerate in the future could be considered much more productive and helpful. If you have this mindset when you decide to complain, you will be less likely to come off as "whiny", and can probably more easily demonstrate that you are simply pointing out a flaw that the person who you are addressing may wish to work on.

ID 4400

The instance above is also annotated as mixed quotation with attitude, since it is made explicit that it is someone else's words. It also used the FLAG since the quoted usage comes from the same conversation.
In Label Studio, this flag is called "SQ FLAG: Previously used in conversation".

- Ambiguous

To be used if there is an ambiguity between two or more SQ labels. It indicates that there is an ambiguity (multiple incompatible readings) between SQ labels.

2.3 Common types of double annotations and ambiguities

In this section we describe and present examples of common types of **double annotations**, i.e., sets of labels that may be annotated together; and **ambiguities**, cases where the context is not sufficient to determine which of a set of labels apply.

As explained above, when doing **double (or triple) annotation,** multiple labels are selected simultaneously because they are considered to apply under the same interpretation of the quoted usage. In **ambiguous** cases, when the context does not allow to determine between two labels which would not simultaneously apply under the same interpretation, the "**Ambiguous**" flag is added.

The list presented here is **not exhaustive**; the annotator may consider that an instance presents a combination of features that is not described here.

If for a QMU the annotator hesitates between two or multiple labels, or sees elements that are compatible with different labels, they need not decide between one of them: they may annotate as many as they think apply to a given QMU.

- Meta-linguistic comment + SQ Usage/Word

Sometimes, a meta-linguistic comment (NSQ) about a word or phrase can express the writer's attitude or opinion toward the word. For example, they may say or imply that, in some way, the word is **problematic, vague, badly defined, inappropriately or unclearly used, or they may point out the word's or usage's ambiguity as a problem**, etc. In the example below, the writer makes a meta-linguistic comment about the fact that the word "meaningful" is not appropriate to characterize what they describe.

Microbrewlogic_##_t1_cnb9dmy_##_rt-t1_cnb9afp: It happens inside me, but I wouldn't say that the nervous signals that govern my bowel movements are "**meaningful**" and that's the same mechanics that govern my "conscious choices" to which there is really only one possible outcome. Why would it be meaningful if there is only one outcome?

io 5585, inner 258

repeatable, observable evidence. Speaking of definitions (hey, you brought it up) I'm not sure *you* understand the definition of "**dogmatic**". You're doing it right now. I'm supposed to take it on faith that you have evidence and yet you repeatedly refuse to present any. If you're making the claim that religion is not dogmatic then yes, you absolutely 100% have to convince me on religion, or at least convince me that there is some sort of evidence that is compelling and does more than pay lip service to the concept of evidence so that people can pretend religion is not dogmatic. Until then you're just being dogmatic. [STA-CITE]> It is not

ID 9176

Since "**both sides**" is a fundamentally inaccurate description of US politics, you have to resort to hyperbole to present a balanced picture. But no balance exists. There just isn't anything on the left that compares to what you are describing on the right. ~~~~~~~~~~~~ TL;DR: I think

ID 287

Such cases constitute a meta-linguistic, typically negative comment about a word, and are annotated with the high-level labels SQ and NSQ, and the fine-grained labels "Meta-linguistic comment" as well as "SQ Usage" (or "SQ Word", as appropriate). The SQ FLAG may also apply.

- Ambiguity: Mixed quotation with/without attitude

It is possible to be uncertain as to whether an instance of mixed quoting expresses attitude (SQ) or not (NSQ). Such cases should be annotated with high-level annotations NSQ, SQ,

“Ambiguous”; and low-level labels “Mixed quotation with attitude” and “Mixed quotation without attitude”.

- SQ usage + Designation

Sometimes, a usage that would usually be considered as a case of designation may come with some attitude.
For example, in some types of designations, quotes serve to indicate that the words go together and are to be interpreted as a unit. However, they may delimit weird or unusual phrasings of concepts, in which case they may simultaneously be considered SQ. These cases are annotated as NSQ, SQ, SQ Usage and Designation.

> autonomy, the ability to improve skills used, and a sense of purpose. The entire purpose of police since their invention is to protect the few who have from the many who have-not. That's all, it's never been about "**protect and serve.**" Just some things to consider.

ID 818

- Meta-linguistic comment + Designation

Some QM can be at the same time meta-linguistic comments and designations. This happens, for example, when someone introduces a specialized term (terminology) and proceeds to define it.

- Meta-linguistic comment + Direct quoting

Sometimes, quotes can serve to indicate that someone said something (direct quotation), and at the same time there is a focus or some meta-linguistic comment on the uttered word. These cases are annotated with NSQ, “Meta-linguistic comment” and “Direct quoting”.

3. Annotation Platform

3.1 Label Studio basic usage

Label Studio is a labeling tool often used for annotation of different phenomena in natural language processing. Annotators will receive a link that will allow them to access the data upon registration. The data in Label Studio is organized in **projects.**

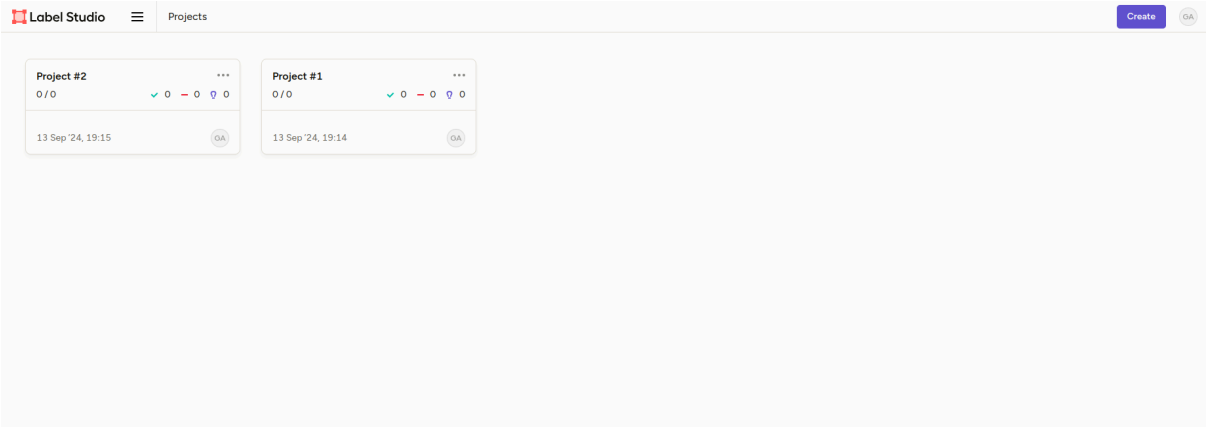

Inside a project, there is one **instance** per line. One instance is focused on a single QMU. Instances can be ordered or filtered with multiple criteria.

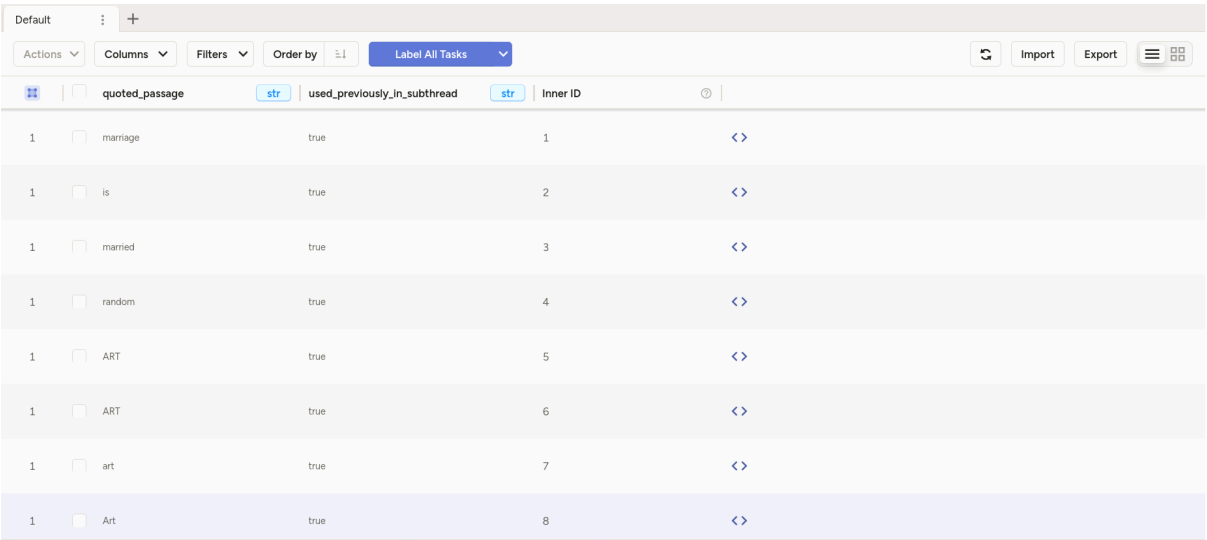

When clicking on one instance, the conversation opens. Every box with text corresponds to a post or the subreddit title. **The QMU for which we need to make a decision is highlighted in blue boldface:**

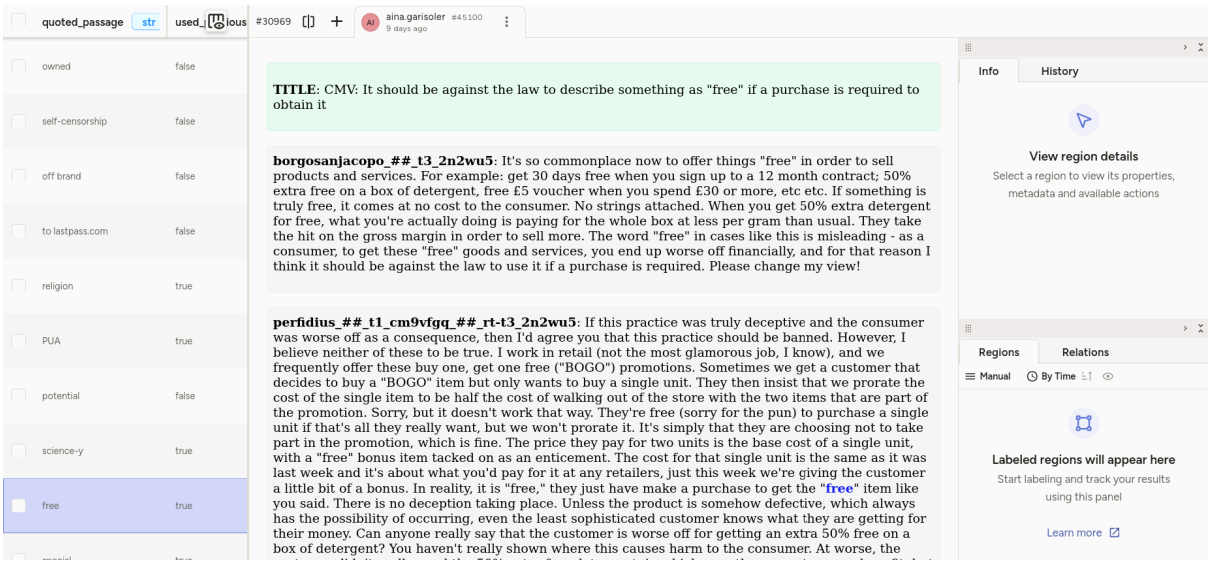

The black boldfaced text at the beginning of each box gives information about the utterance in this format:
`[USERNAME]_##_[POST ID]_##_[POST ID REPLIED TO]`
You will probably only need to use the username, which comes before the first underscore. For example, in the screenshot above, the QMU to be annotated has been written by a user with username “perfidius”.

The conversation shown in an instance does not contain all posts in it: it includes only those belonging to the subthread that involves the QMU (the chain of posts it replies to, and the posts that reply to it).

Scrolling down to the end of the conversation, one can find the options to select to annotate an instance. Once annotation is done, it needs to be submitted by clicking on “Submit” or by pressing Ctrl+Enter. Moving to the next instance can be done more quickly by pressing Shift+down arrow.

More information on how to customize columns and apply filters will be given in an interactive session with the author of these guidelines.

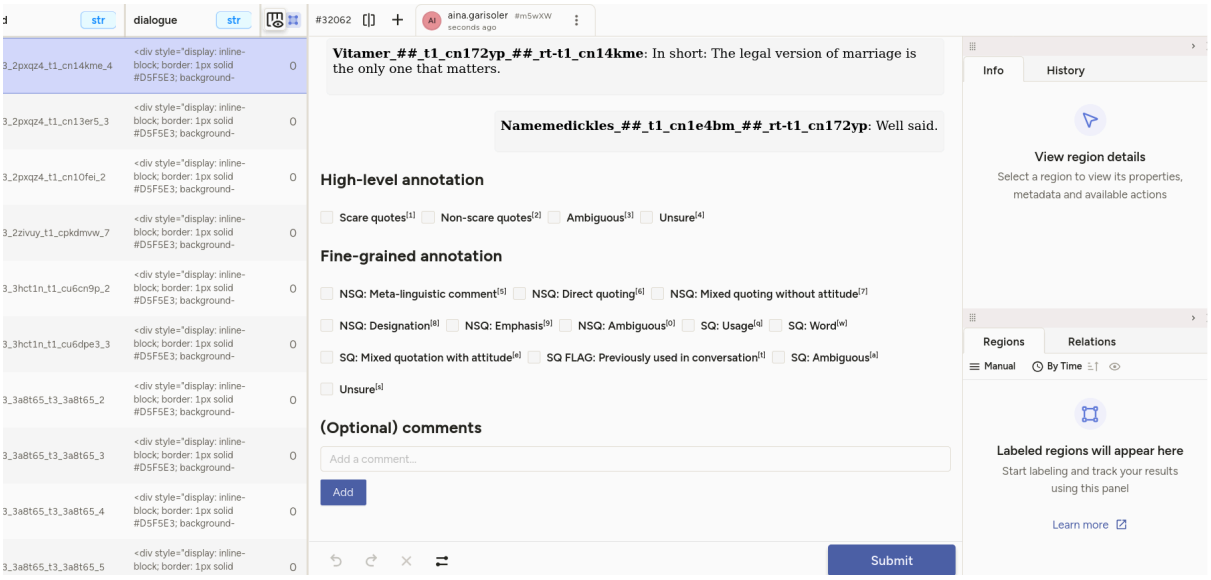

3.2 Reddit and format specificities

The corpus we use, winning-args, dates from 2013-2015.

Quotations

Redditors may quote other users' posts or external information. These passages are delimited with the signs `[STA-CITE]` and `[END-CITE]`, with very few exceptions due to formatting irregularities in the original corpus. To avoid redundancy, in this study we have ignored all QMUs that take place between these two signs.

Here is an example:

> **johnpseudo_##_t1_cbqvavl_##_rt-t1_cbqv7cr**: [STA-CITE]>why do you see the only result being the end civilization? [END-CITE]I don't understand your question. [STA-CITE]>it seems unlikely that all of humanity could not come up with any project to save ourselves. [END-CITE]Why does that seem unlikely? It's happening very slowly right now. [STA-CITE]>if humanity can accidentally make the world warmer, why couldn't we intentionally make the world cooler? [END-CITE]Because of inertia and fundamental limitations of physics. We already had the knowledge and technology to address this years ago, but we continue to increase our carbon emissions. Once we finally reach a point where the consensus favors action, we will have already reached a point where the technology to reverse the change is beyond our reach.

Edited messages

Sometimes redditors edit their messages. They often add something about their edit at the end of the post: they may include the reason for their edit or additional text. It is left to the judgment of the annotator to decide whether an edit should be part of a WMN or not.

Here is an example:

> isn't particularly convincing to me. So...change my view? Is there a better argument that I'm not considering here? I'm interested in opinions both from a moral and a practical/legal perspective. (edit: ack formatting!)

Reddit / winning-args corpus terminology

- **OP**: Original Poster (the user who started the thread)
- **CMV:** abbreviation of the subreddit name "changemyview".
- **Deltas:** The winning-args corpus contains posts from the r/changemyview subreddit, where the OP presents their opinion on an issue and other users challenge it or give different perspectives. When a user manages to change the OP's opinion in some respect, the OP "awards a delta" to that user and includes a delta symbol in their reply (Δ).

4. Simplified action plan

1. Find the QMU in blue boldface and read the sentence containing it.
2. Read the surrounding text in the given post, and preceding posts if you need more context.
3. High-level decision: is it SQ or NSQ, both, or is it ambiguous?
4. Low-level decision:
    a. If NSQ: select one or more NSQ labels that apply.
    b. If SQ: select one or more SQ labels that apply. Verify if you need to add the FLAG.
    c. if both SQ and NSQ: select one or more NSQ and SQ categories (make sure to annotate at least one of each). Verify if you need to add the FLAG.
5. If you annotated multiple phenomena and they are not compatible, make sure to mark it as “Ambiguous” at the concerned level.
6. If you are unsure of your annotation, mark it as “Unsure”.